


Introduction to Cilk

Programming Models for Emerging Platforms

First Model: POSIX Threads

- C library for thread programming
- Almost as crude as it gets



What does this mean?

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

“A bad day writing code in Scheme is better than a good day writing code in C. “

- David Stigant

“A bad day writing code in Scheme is better than a good day writing code in C. “

- David Stigant

My point: There is more to world than C and “performance”

Programming Languages

- We research programming languages to:
 - Provide stronger guarantees (garbage collection = memory safety)
 - Make writing software more productive (spending hours / days / months searching for memory / concurrency bugs is not)
 - Make writing software more *enjoyable*
- A programming model is a subset of programming languages

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 -
 - It's unlikely we remove all items on this list,
But we can at least shrink it
 -
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

First Model: POSIX Threads

- With threads, programmers must:
 - 1. Manage physical unit of execution (thread)
 - 2. Manage logical unit of execution (code)
 - 3. Synchronize these physical and logical units of execution
 - 4. Handle life cycle dependency *between* threads
 - 5. Handle memory dependency *between* threads
 - 6. Deal with challenging side effect of concurrency: race conditions and deadlock

Second Model: Cilk

- C language extension for writing parallel and concurrent programs
- Developed at MIT in late 90s
- Part of gcc toolchain (search for intel cilkplus)
- No more threads! We use workers (which are invisible to the programmer) and spawn work.

Second Model: Cilk

- Work through example (fib.c)

Second Model: Cilk

- Key to Cilk is the idea of *work stealing*...

1



2




3




4







1



2



3




4







1



2




3



4









1




Victim



2



3




4

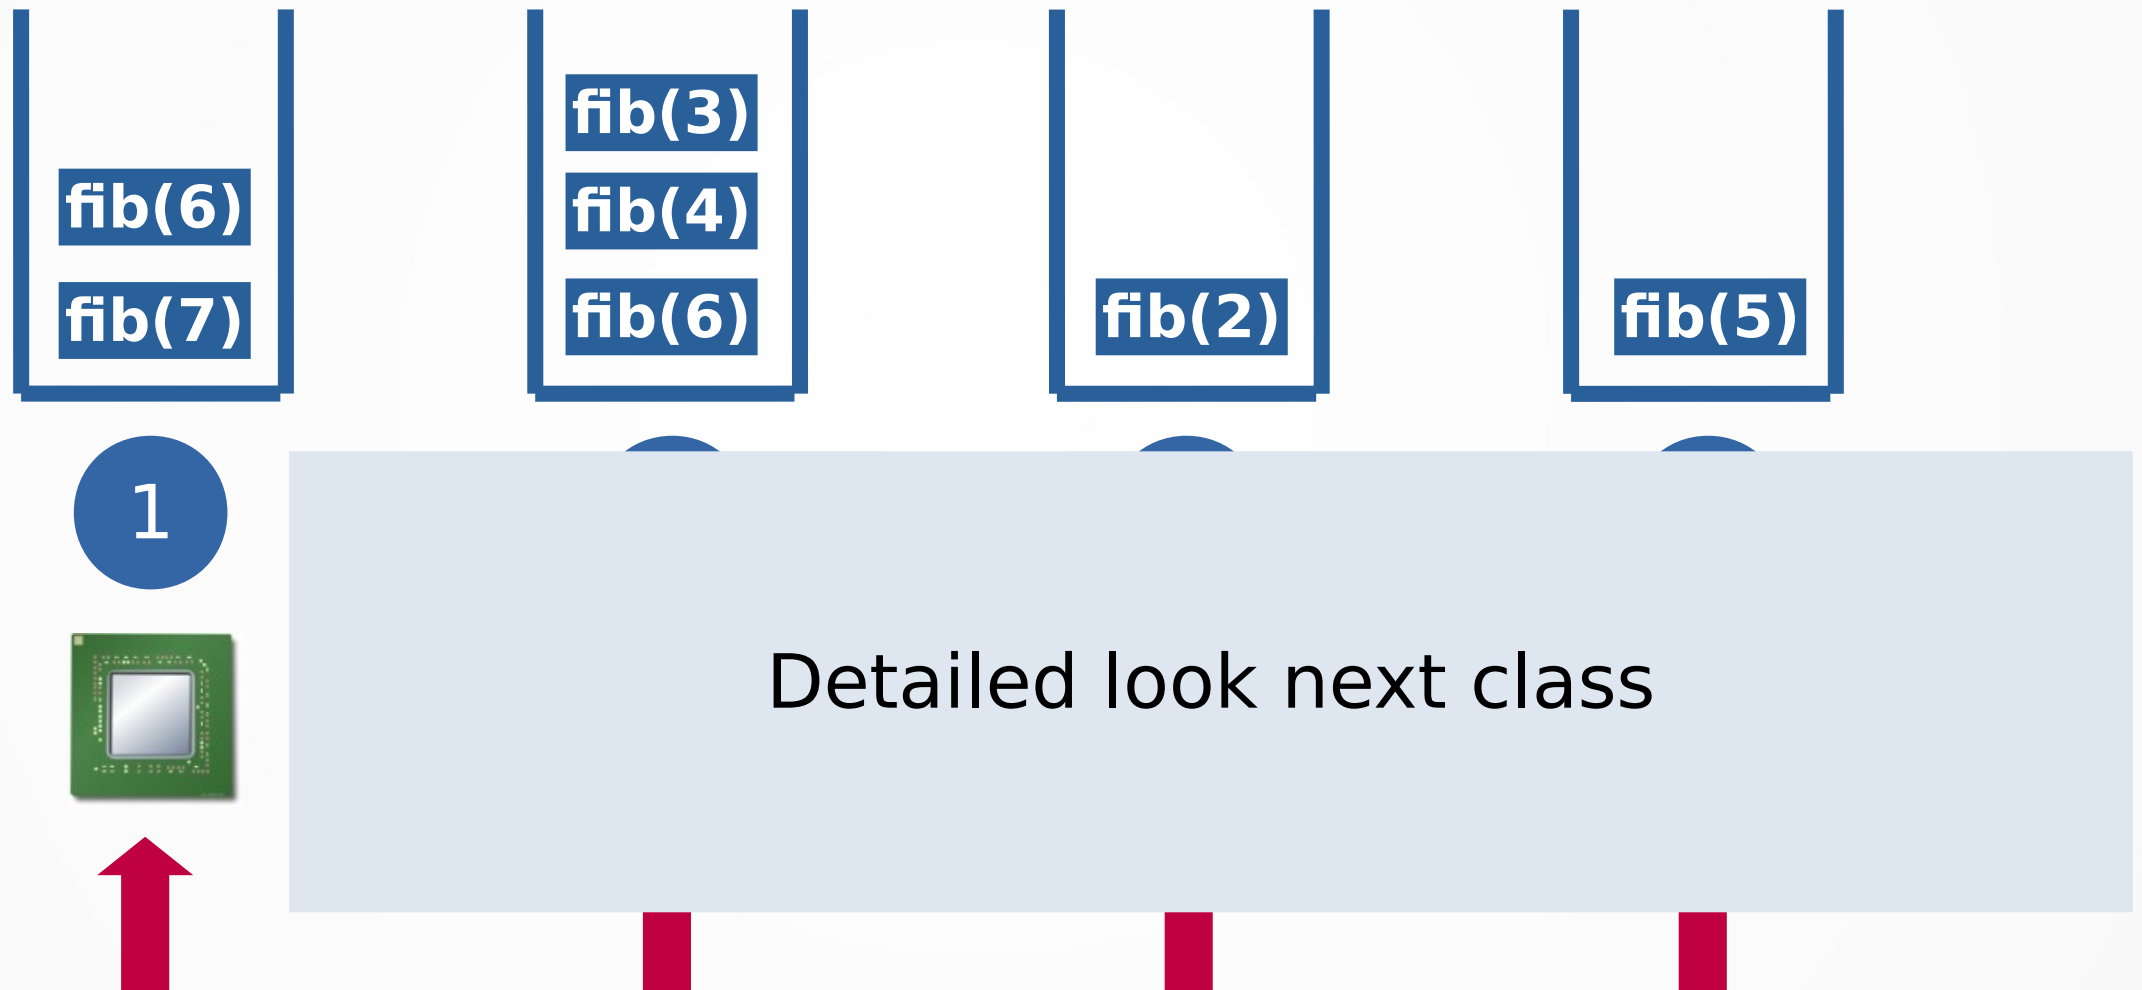


Thief

Dynamic Load Balancing



Dynamic Load Balancing



Second Model: Cilk

- Work through example (vector.c, merge.c)

For Reference / Self Study

- cilk statements
 - `cilk_spawn`
 - `cilk_sync`
 - `cilk_for`
- Compiling and linking
 - Use cilk header: `#include <cilk/cilk.h>`
 - Compile: `gcc -std=c11 -fcilkplus code.c`