

MPI Design and Implementation

Programming Models for Emerging Platforms


MPI *Standard*



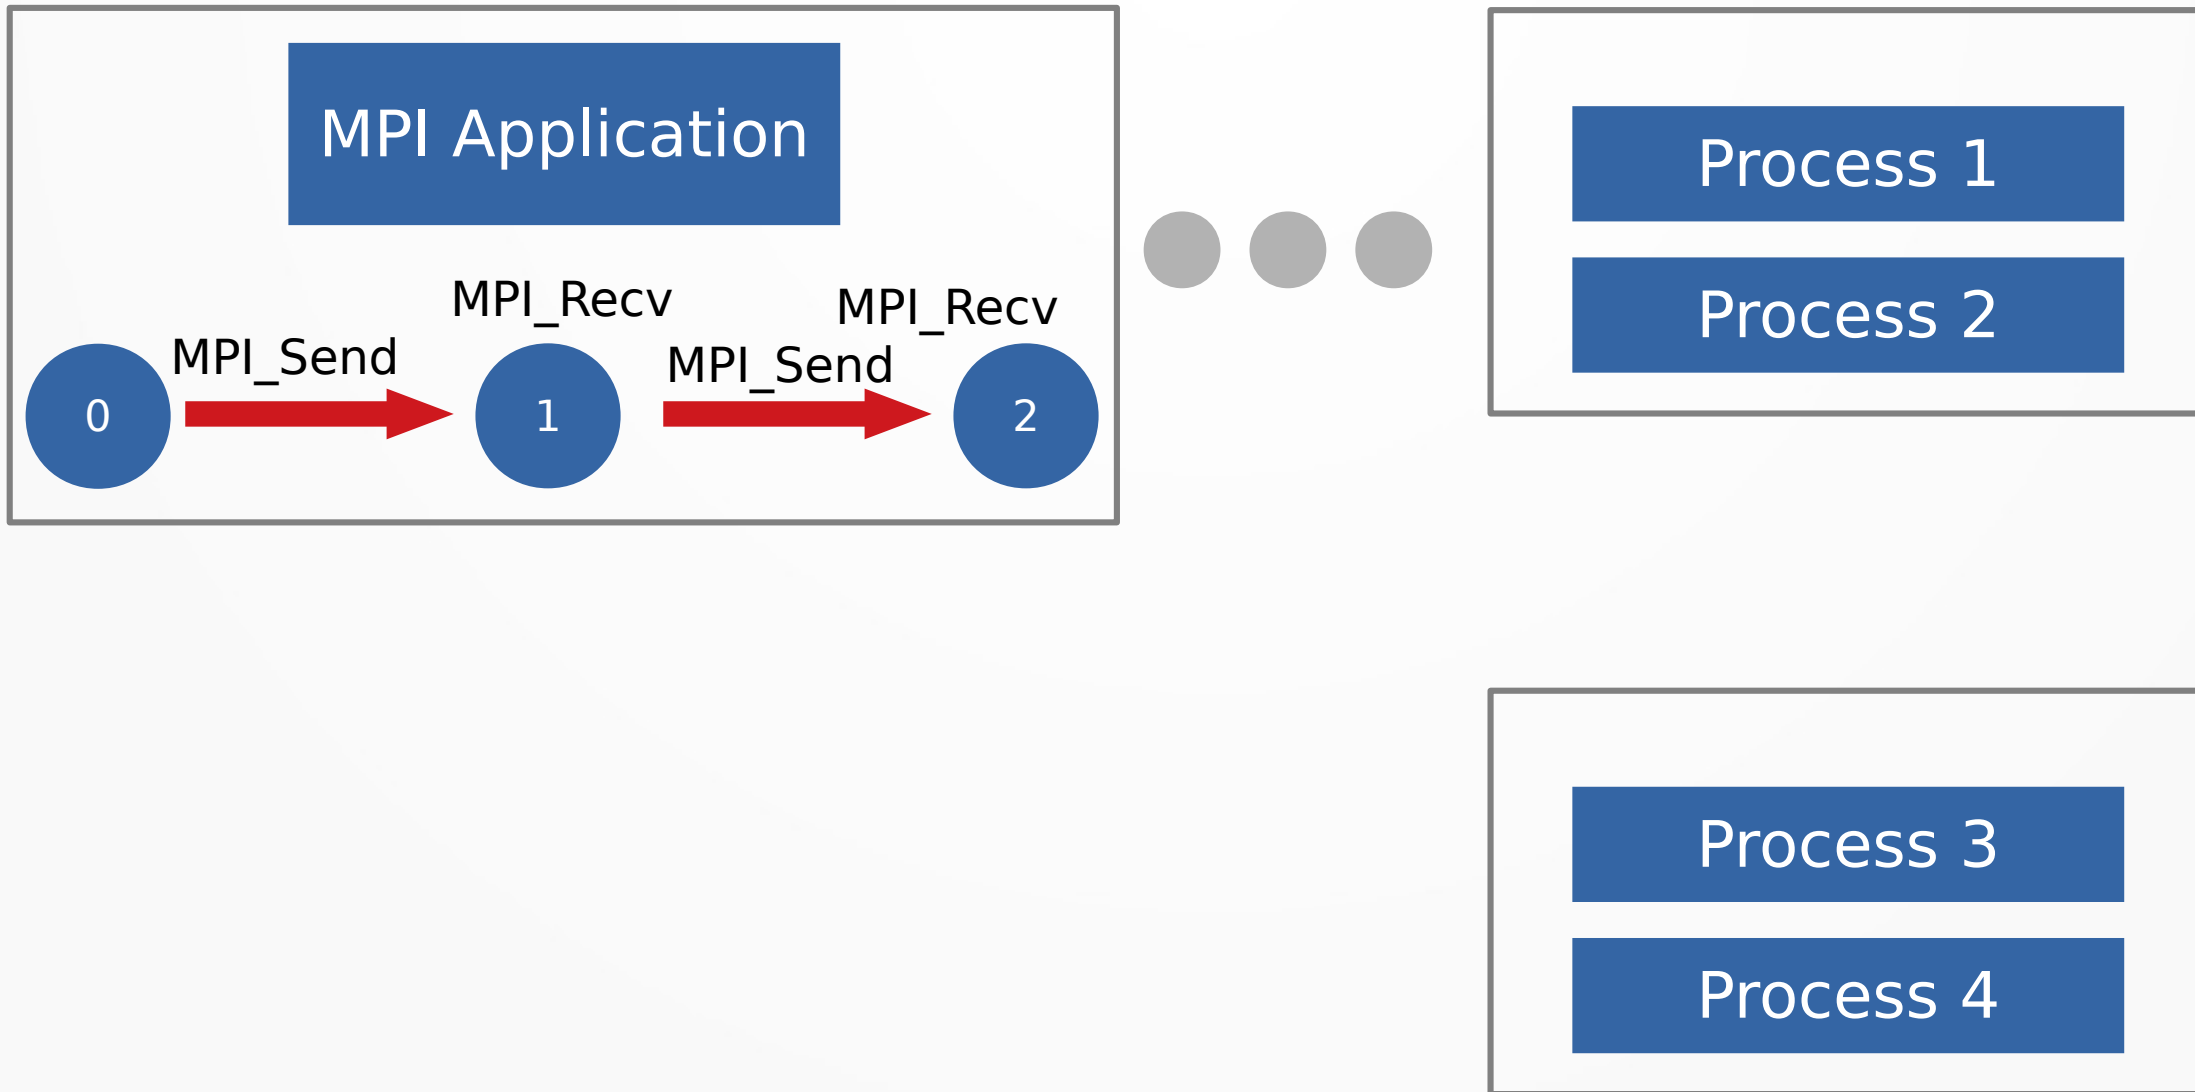
MPI Application

MPI *Standard*


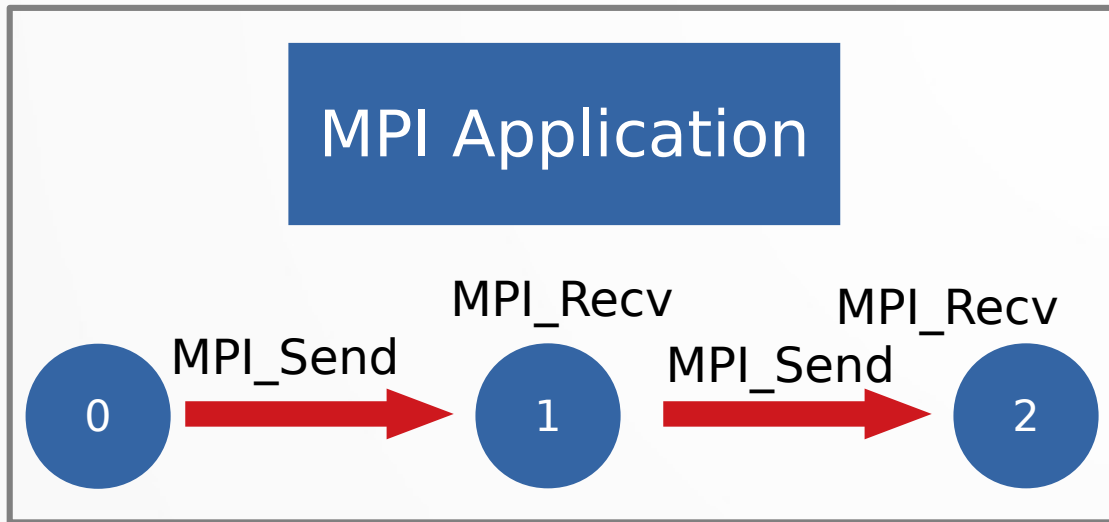
MPI Application




MPI *Standard*



MPI *Standard*

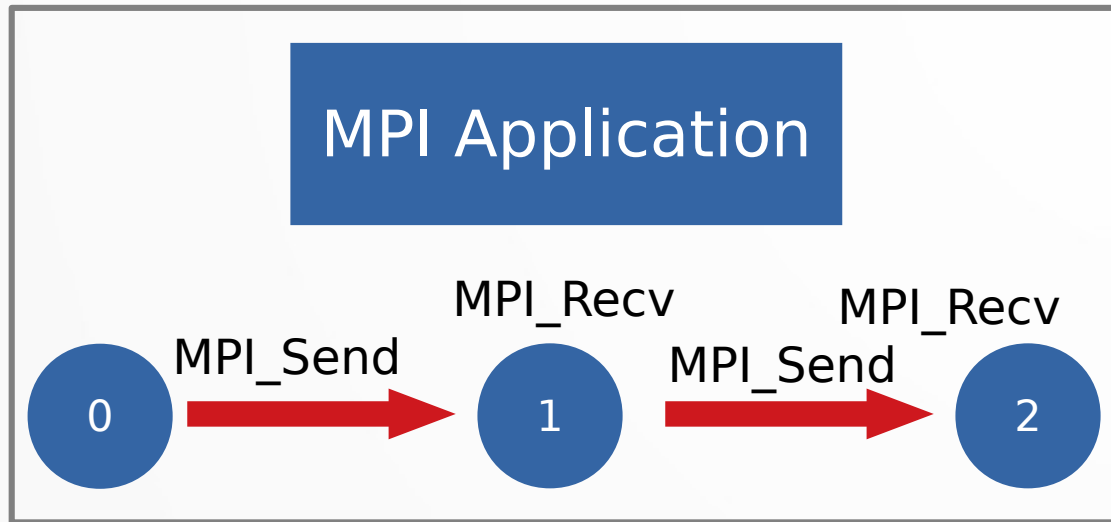


Communicate over
High-Performance Switches

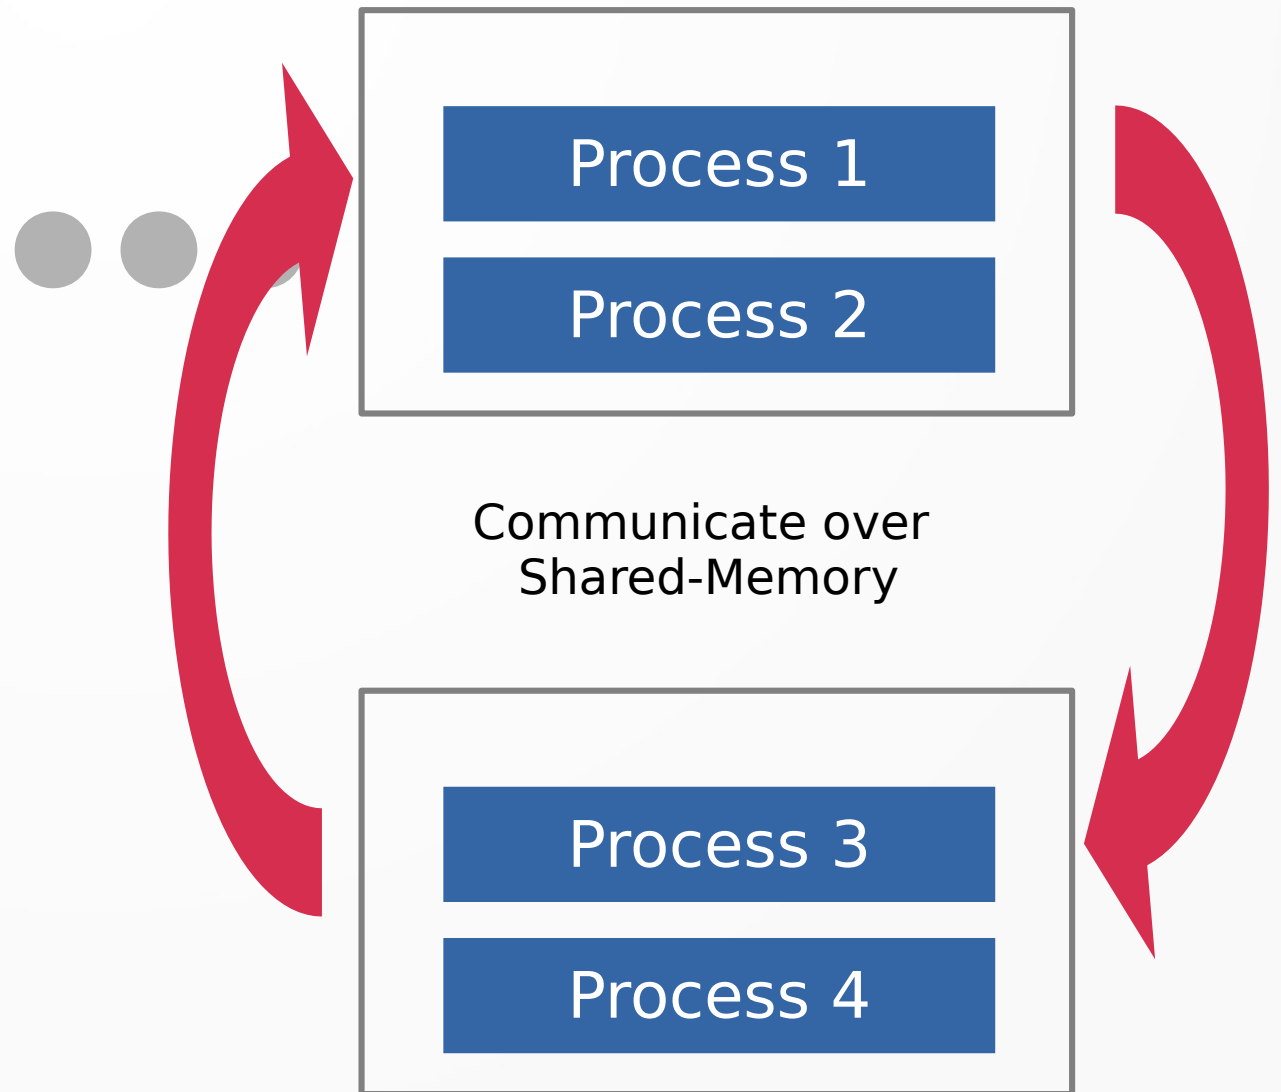


**Distributed Memory Parallel
Supercomputer**


MPI *Standard*



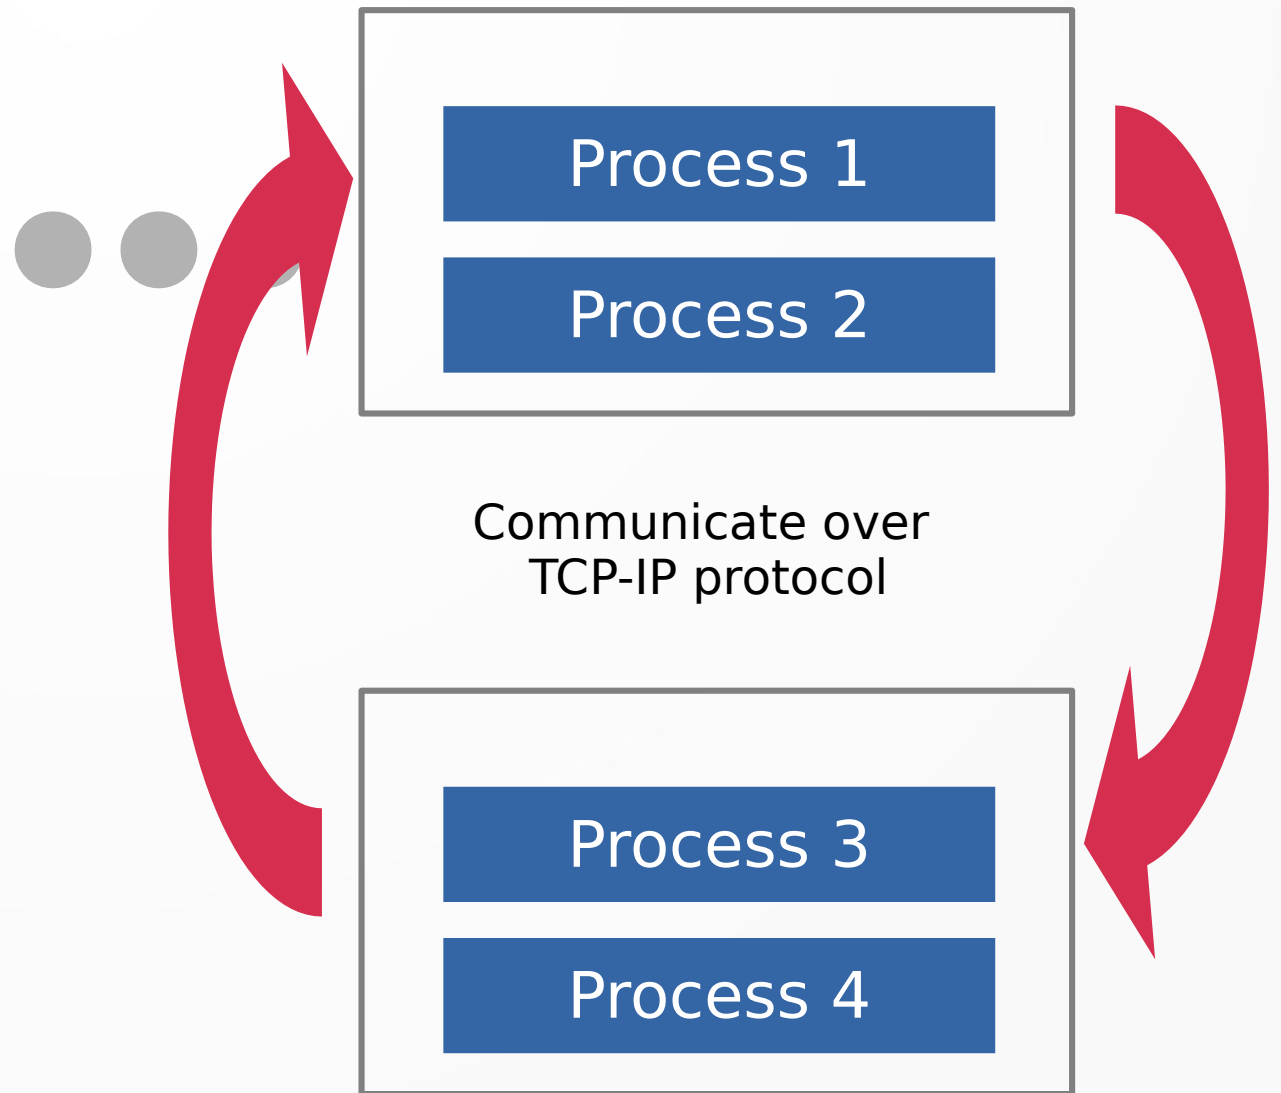
Shared-Memory Architecture



MPI *Standard*



Network of Workstations



MPI *Standard*

Standard (Abstraction Layer) allows programmer to elide implementation details

Network of Workstations

Fill in

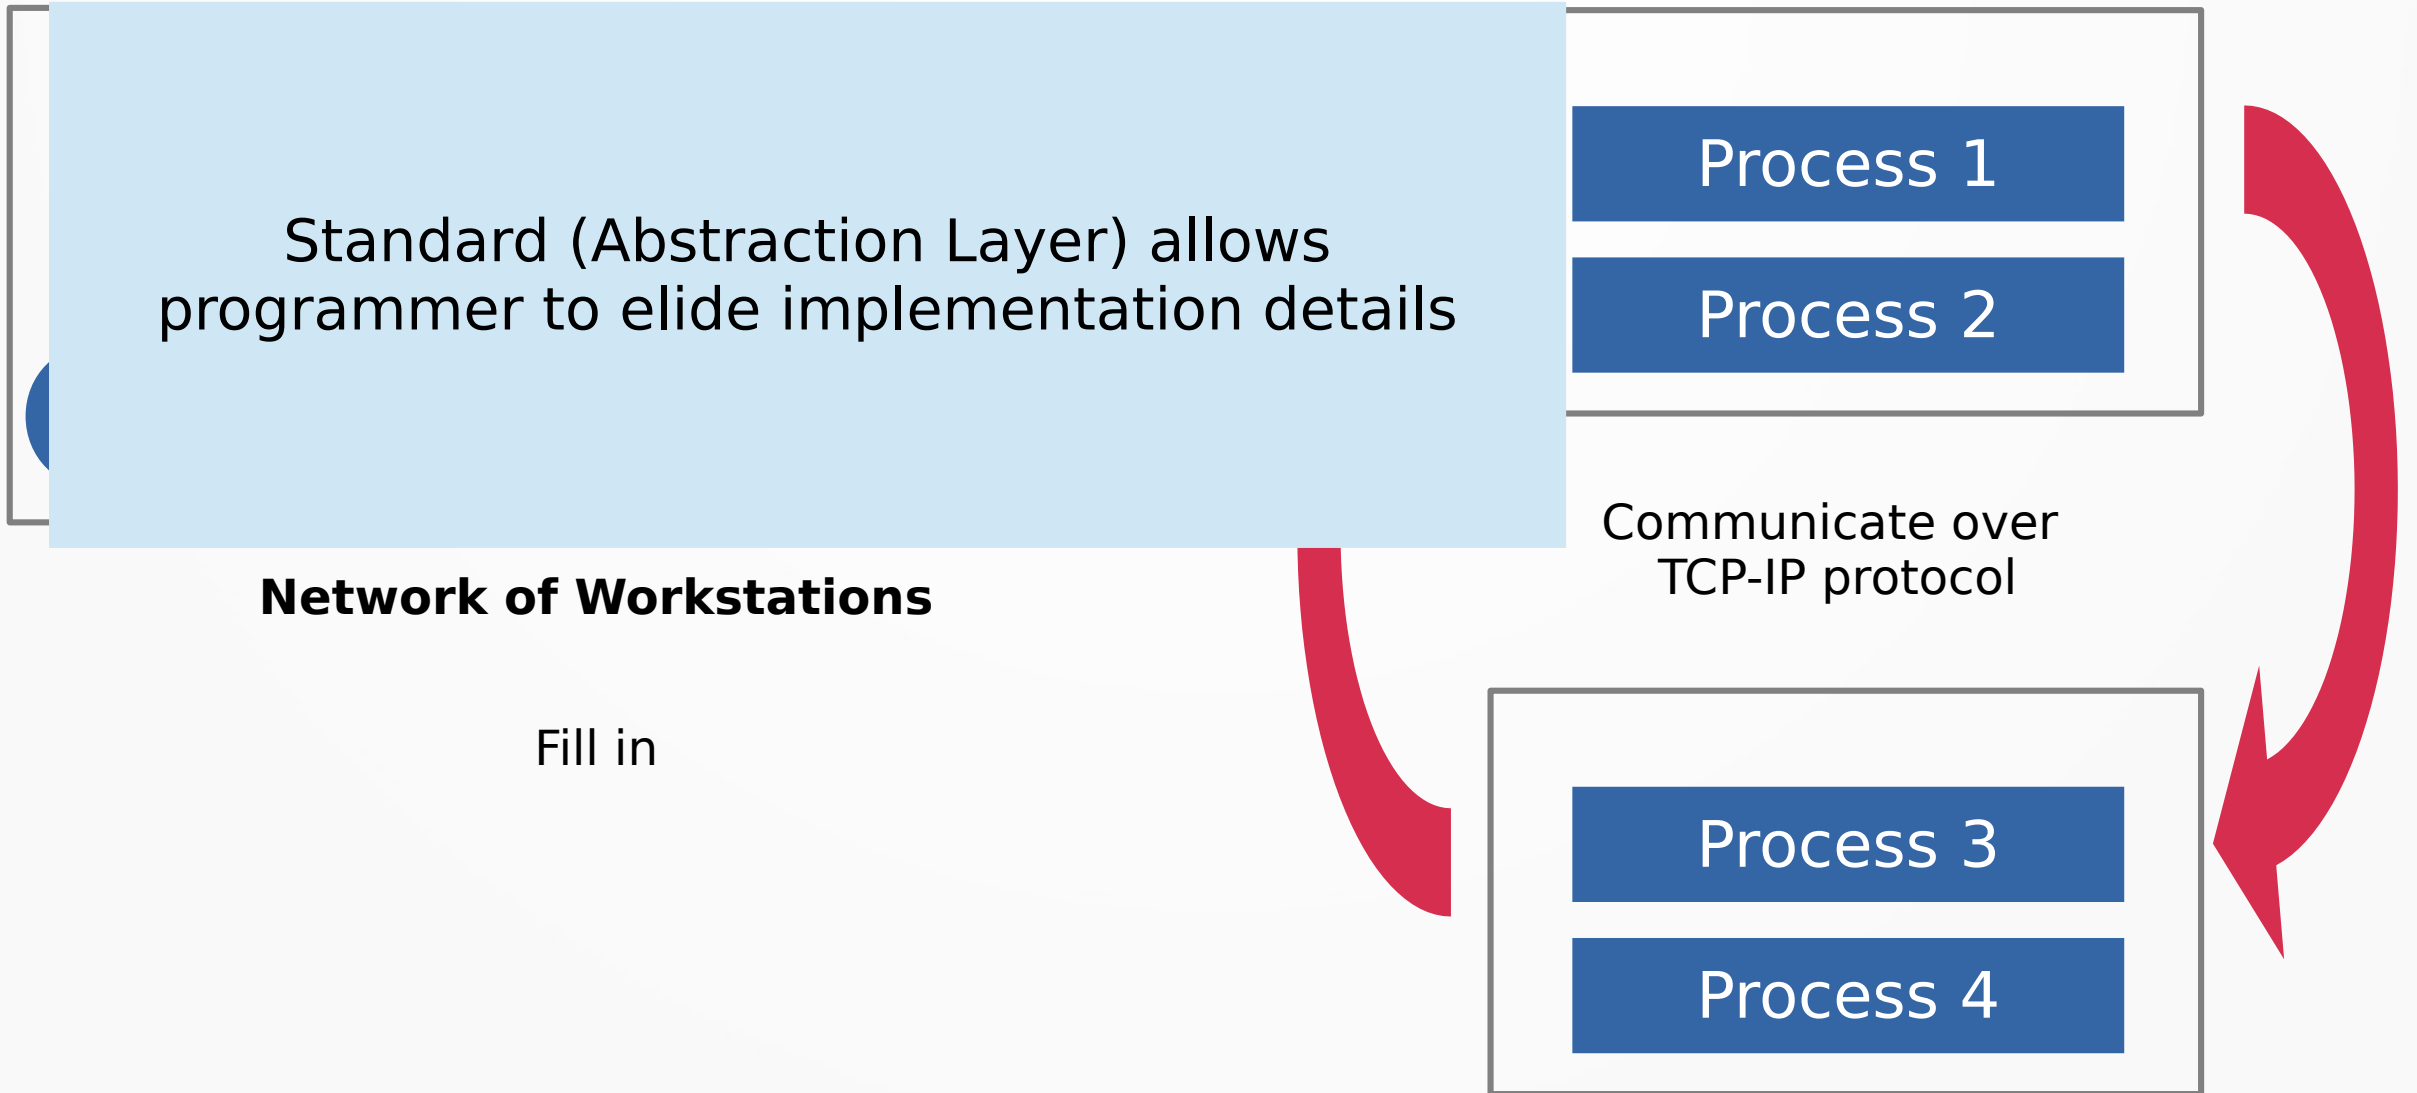
Process 1

Process 2

Communicate over
TCP-IP protocol

Process 3

Process 4



MPI *Standard*

Standard (Abstraction Layer) allows programmer to elide implementation details

Process 1

Process 2

Communicate over
TCP/IP

Network of Works


Fill in

Write once, run anywhere
(Java didn't invent this)



Compares MPICH with native machine communication interface (NX)


MPICH uses NX on the Paragon machine




Compares MPICH with native machine communication interface (NX)

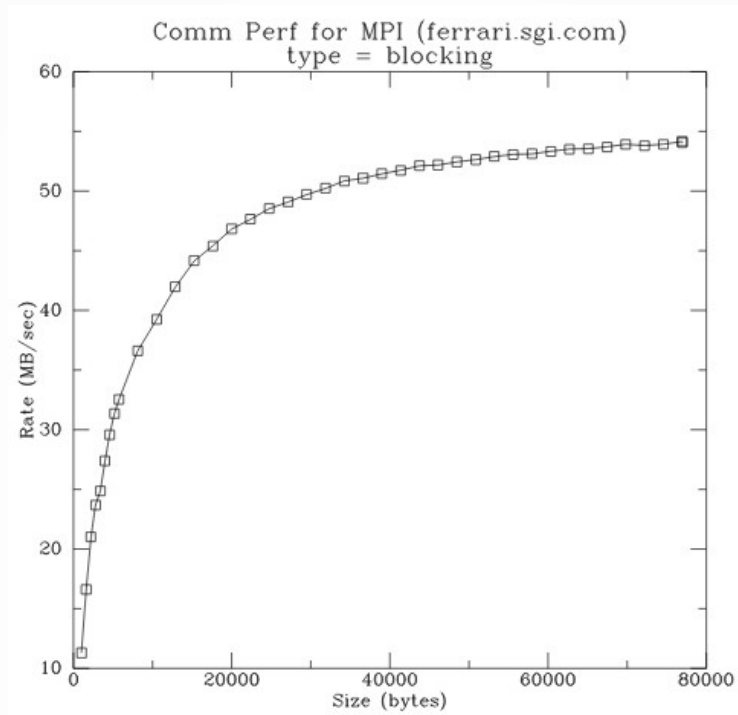
MPICH uses NX on the Paragon machine

There is almost always a *cost* of abstraction

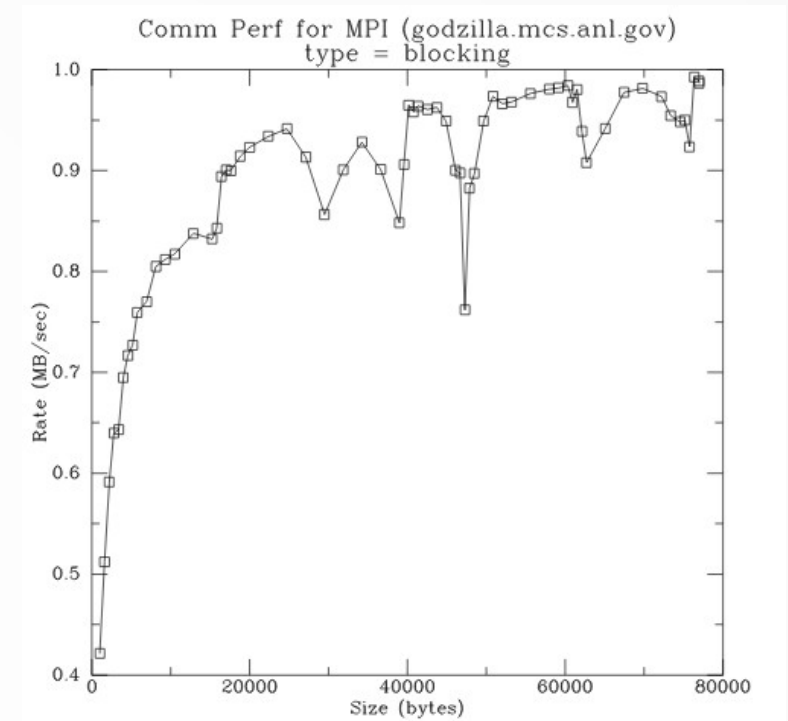




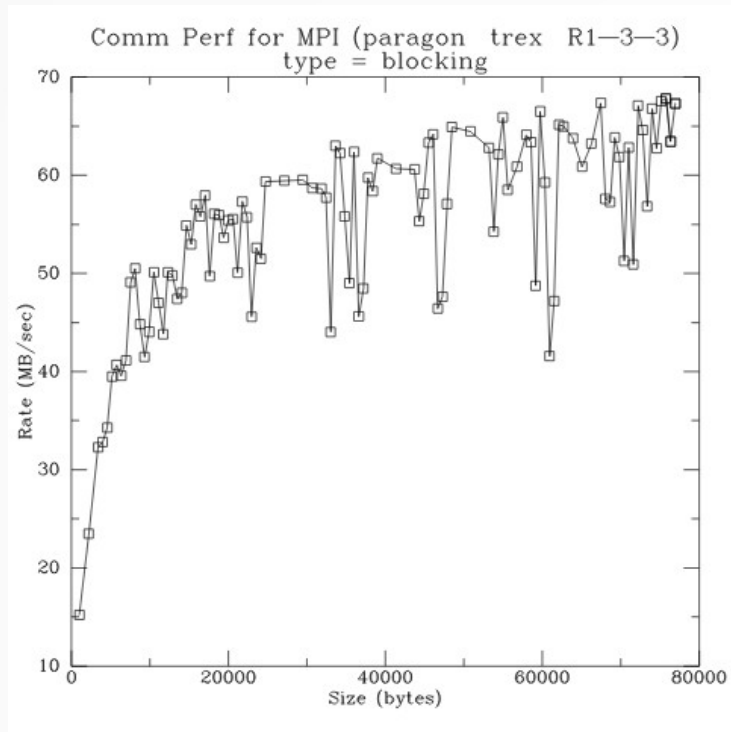
High-Performance
Switch



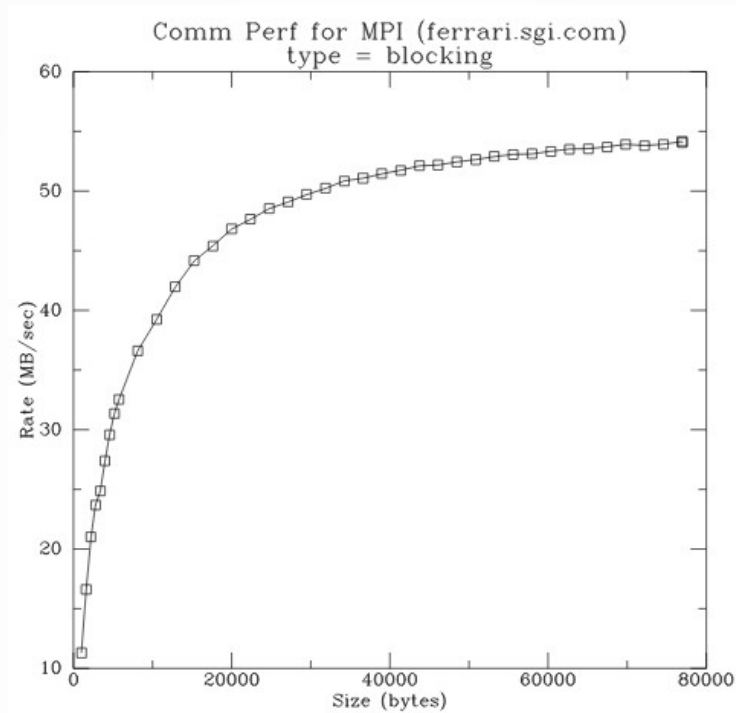
Shared-Memory




Workstation Network



High-Performance
Switch



Shared-Memory



Workstation Network

View this as increasing the range of possible programming environments, not which machine / environment is best

How to *implement* the standard to be portable and extensible?

Interface Design

- We often say “provide an interface” as if it is a simple, straightforward part of API design

Interface Design

- We often say “provide an interface” as if it is a simple, straightforward part of API design
- It’s not! Particularly when lower-layers are involved

Interface Design

- We often say “provide an interface” as if it is a simple, straightforward part of API design
- It’s not! Particularly when lower-layers are involved
- If you expose too little to the interface, can’t do anything useful. If you expose too much, might as well not have an interface because you must re-implement everything

Interface Design

- Consider process manipulation (MPI needs processes)
- What even is a “process”? Is a “process” the same on all machines? Specialized hardware? Across network?
- All this needs abstraction

Abstract Device Interface

- As far as MPI goes, a device needs to provide functions for:
 - 1. Specifying a message to be sent or received
 - 2. Moving data between API and message-passing hardware (User space, device space)
 - 3. Manage pending messages (send and receive)
 - 4. Provide information about execution environment

MPI_Reduce

MPI

MPI_Isend

**MPI
point-to-point**

MPID_Post_Send

**The Abstract
Device Interface**

SGI(4)

MPID_SendControl

**The Channel
Interface**

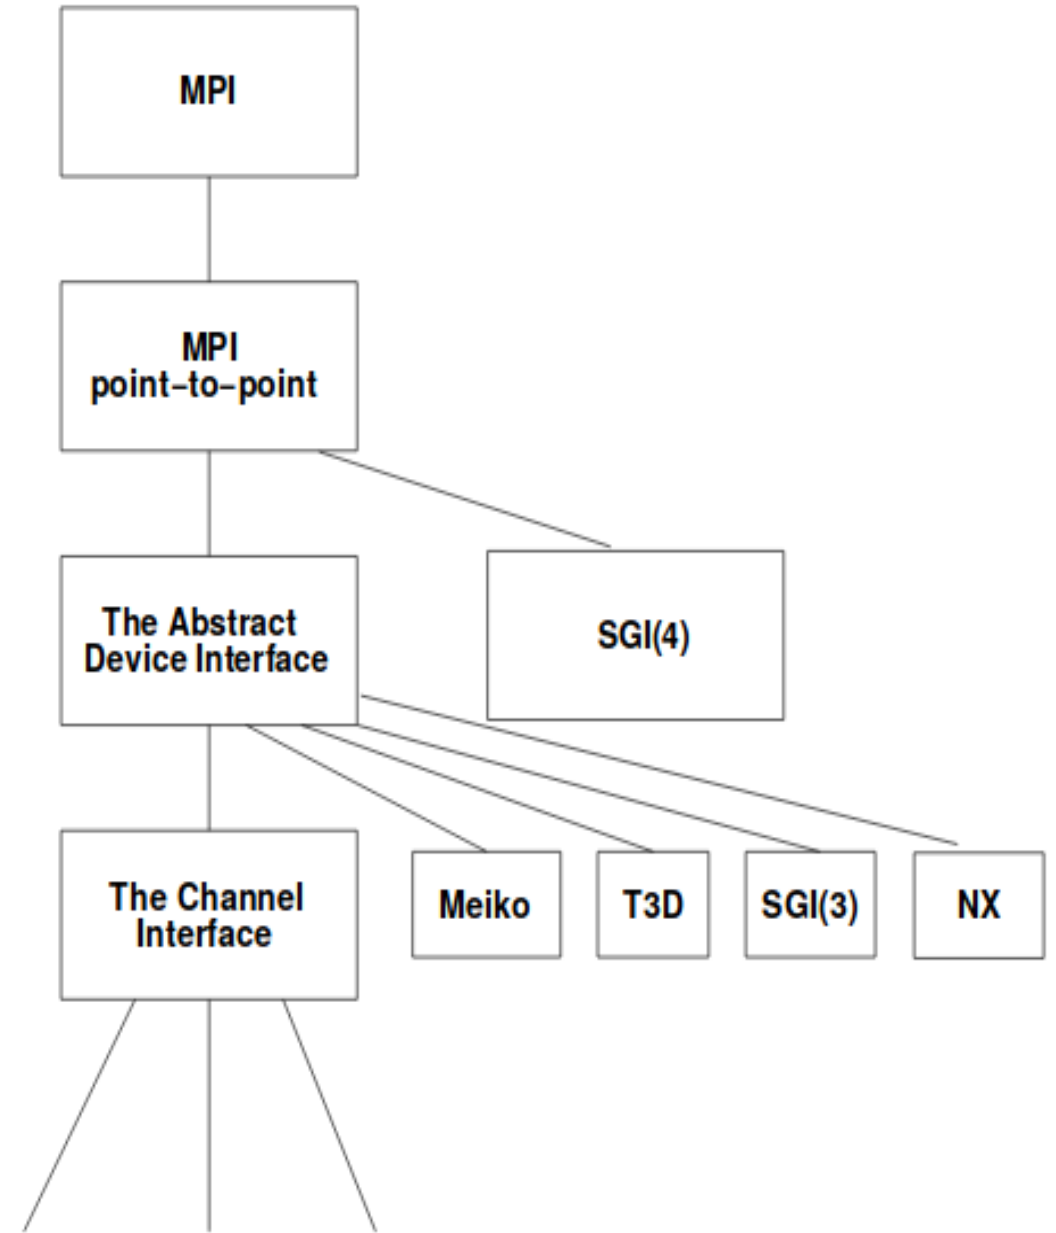
Meiko

T3D

SGI(3)

NX

(implementations of the channel interface)



Collective communication
implemented over
point-to-point

Send moves from API to
(ADI) referred to as MPID in
code

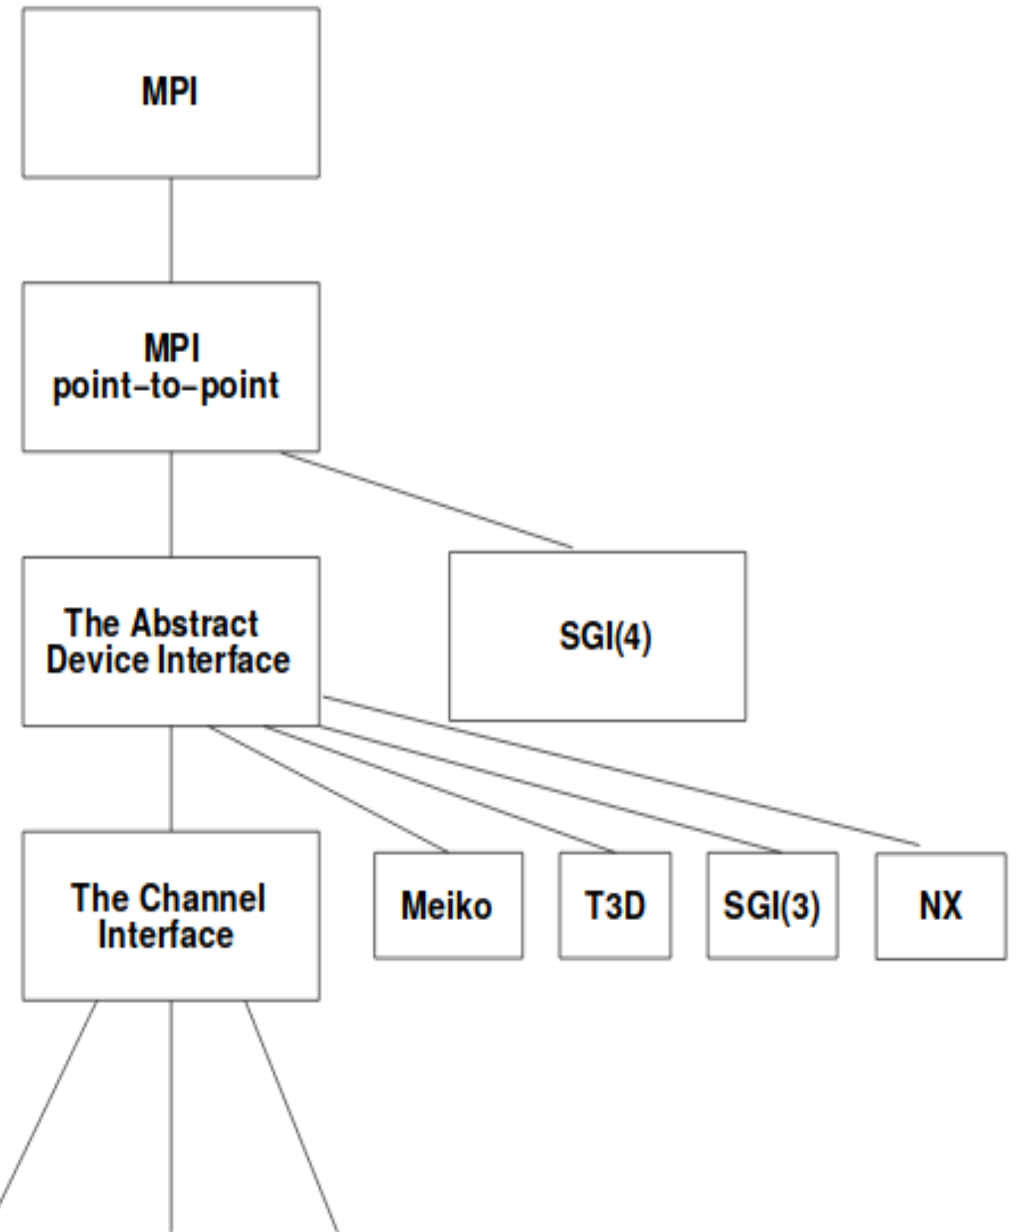
MPI_Reduce

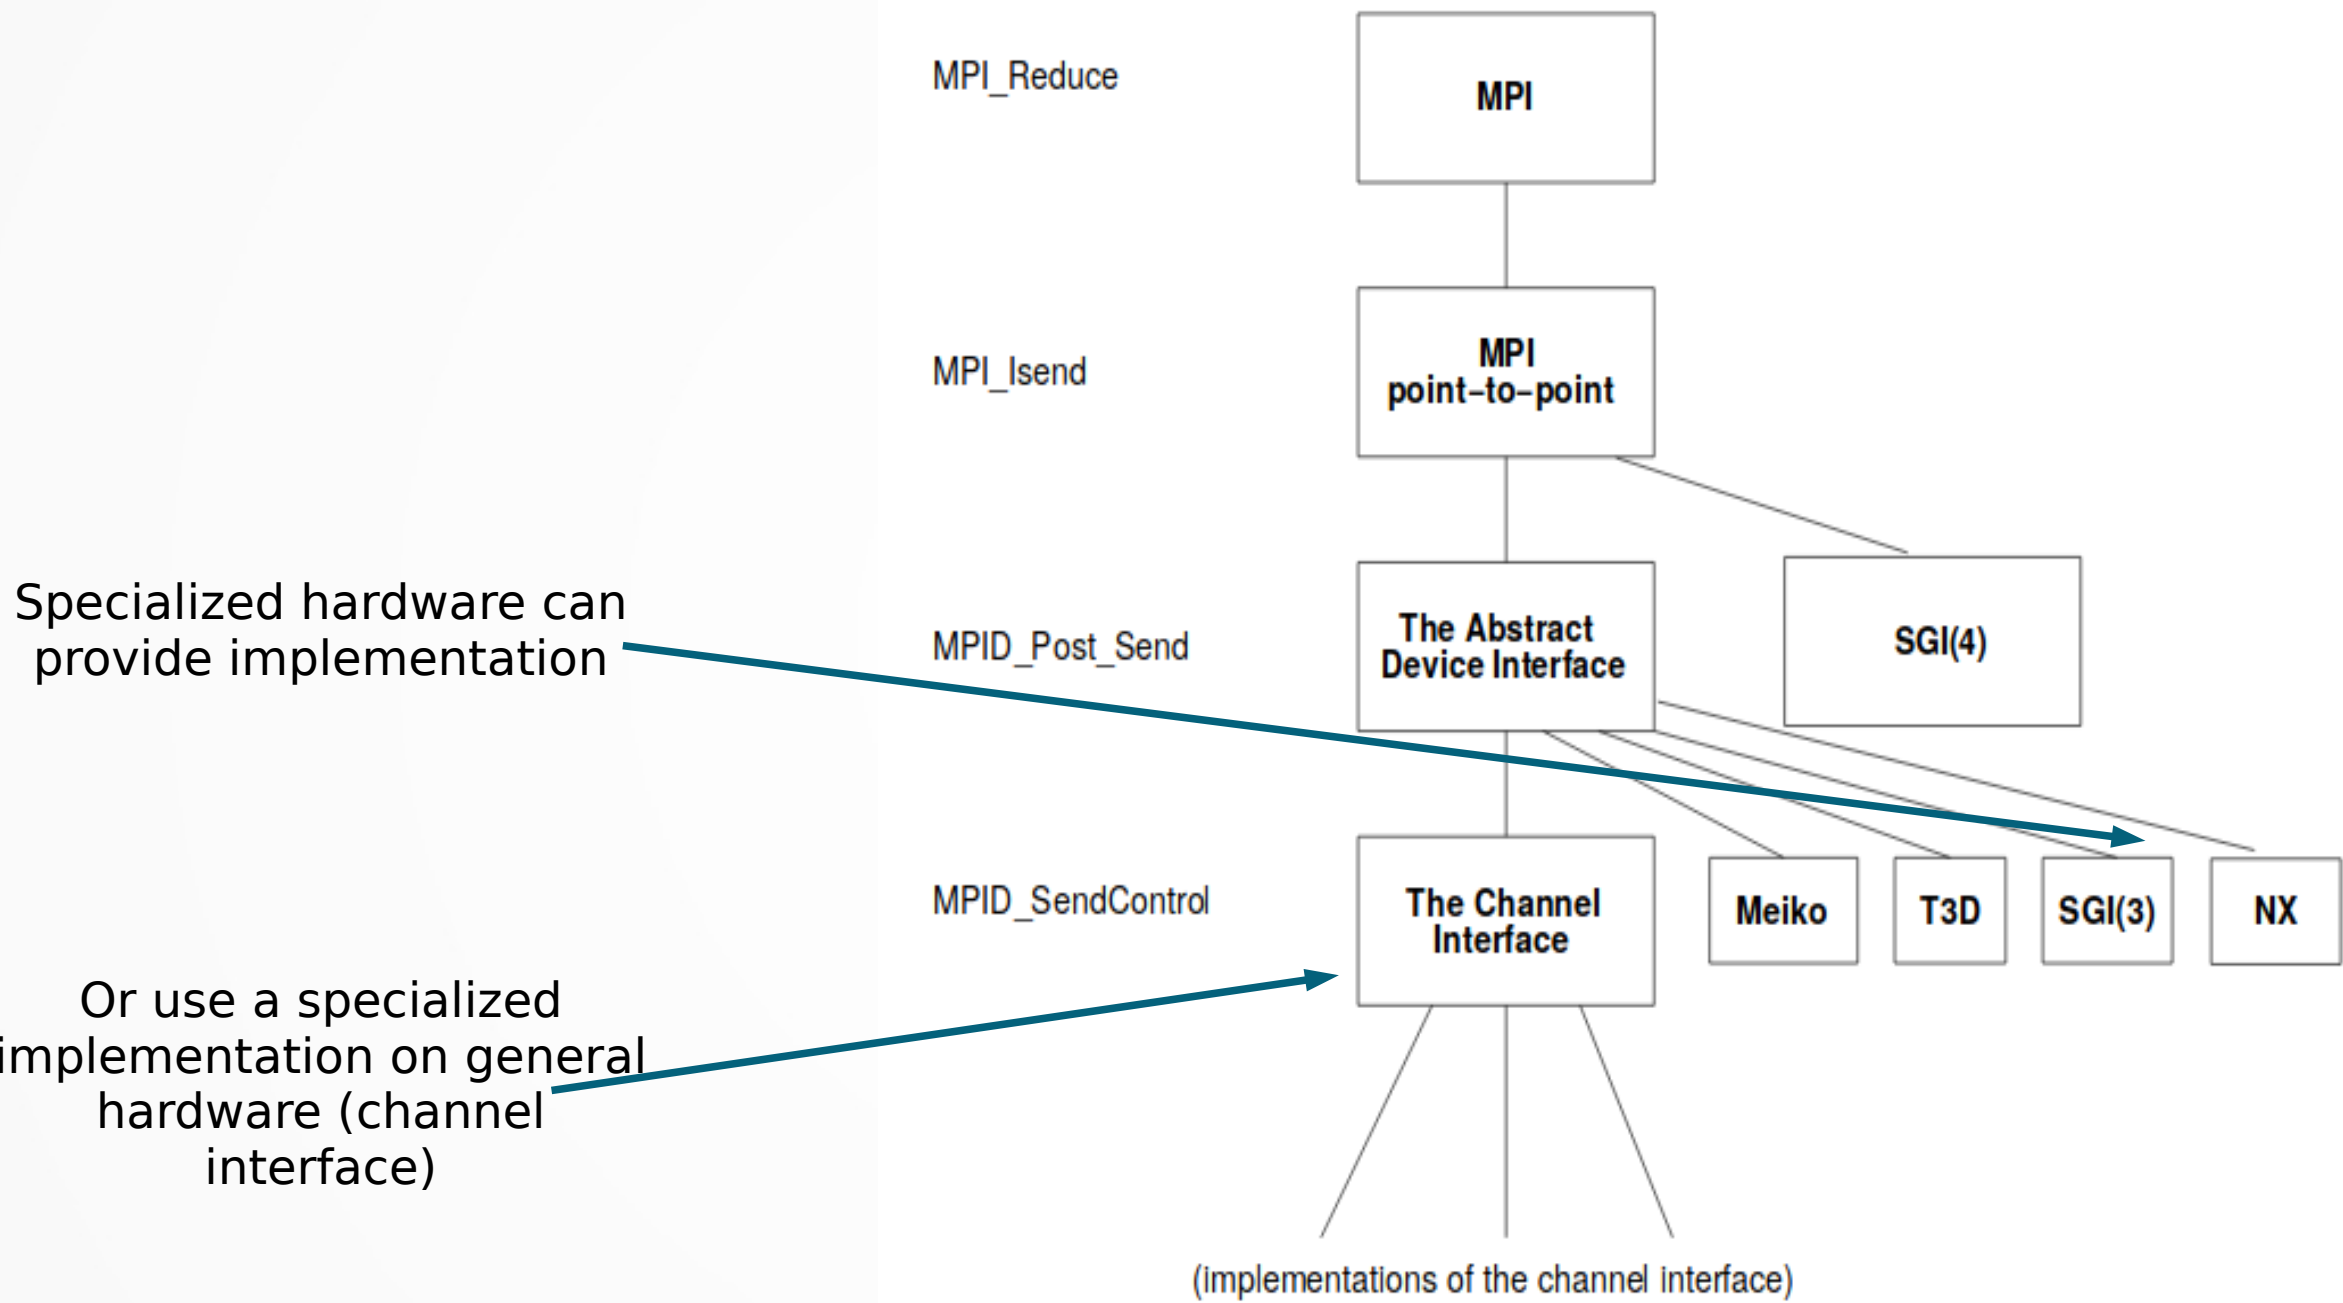
MPI_Isend

MPID_Post_Send

MPID_SendControl

(implementations of the channel interface)





Channel Interface

- Low-level functions to physically transfer data from one process space to another process space
- Remember, MPI abstracted lots of details. Even sockets abstract details, but its not that simple...
- Data exchange mechanisms
- Eager
- Rendezvous
- Get

Channel Interface

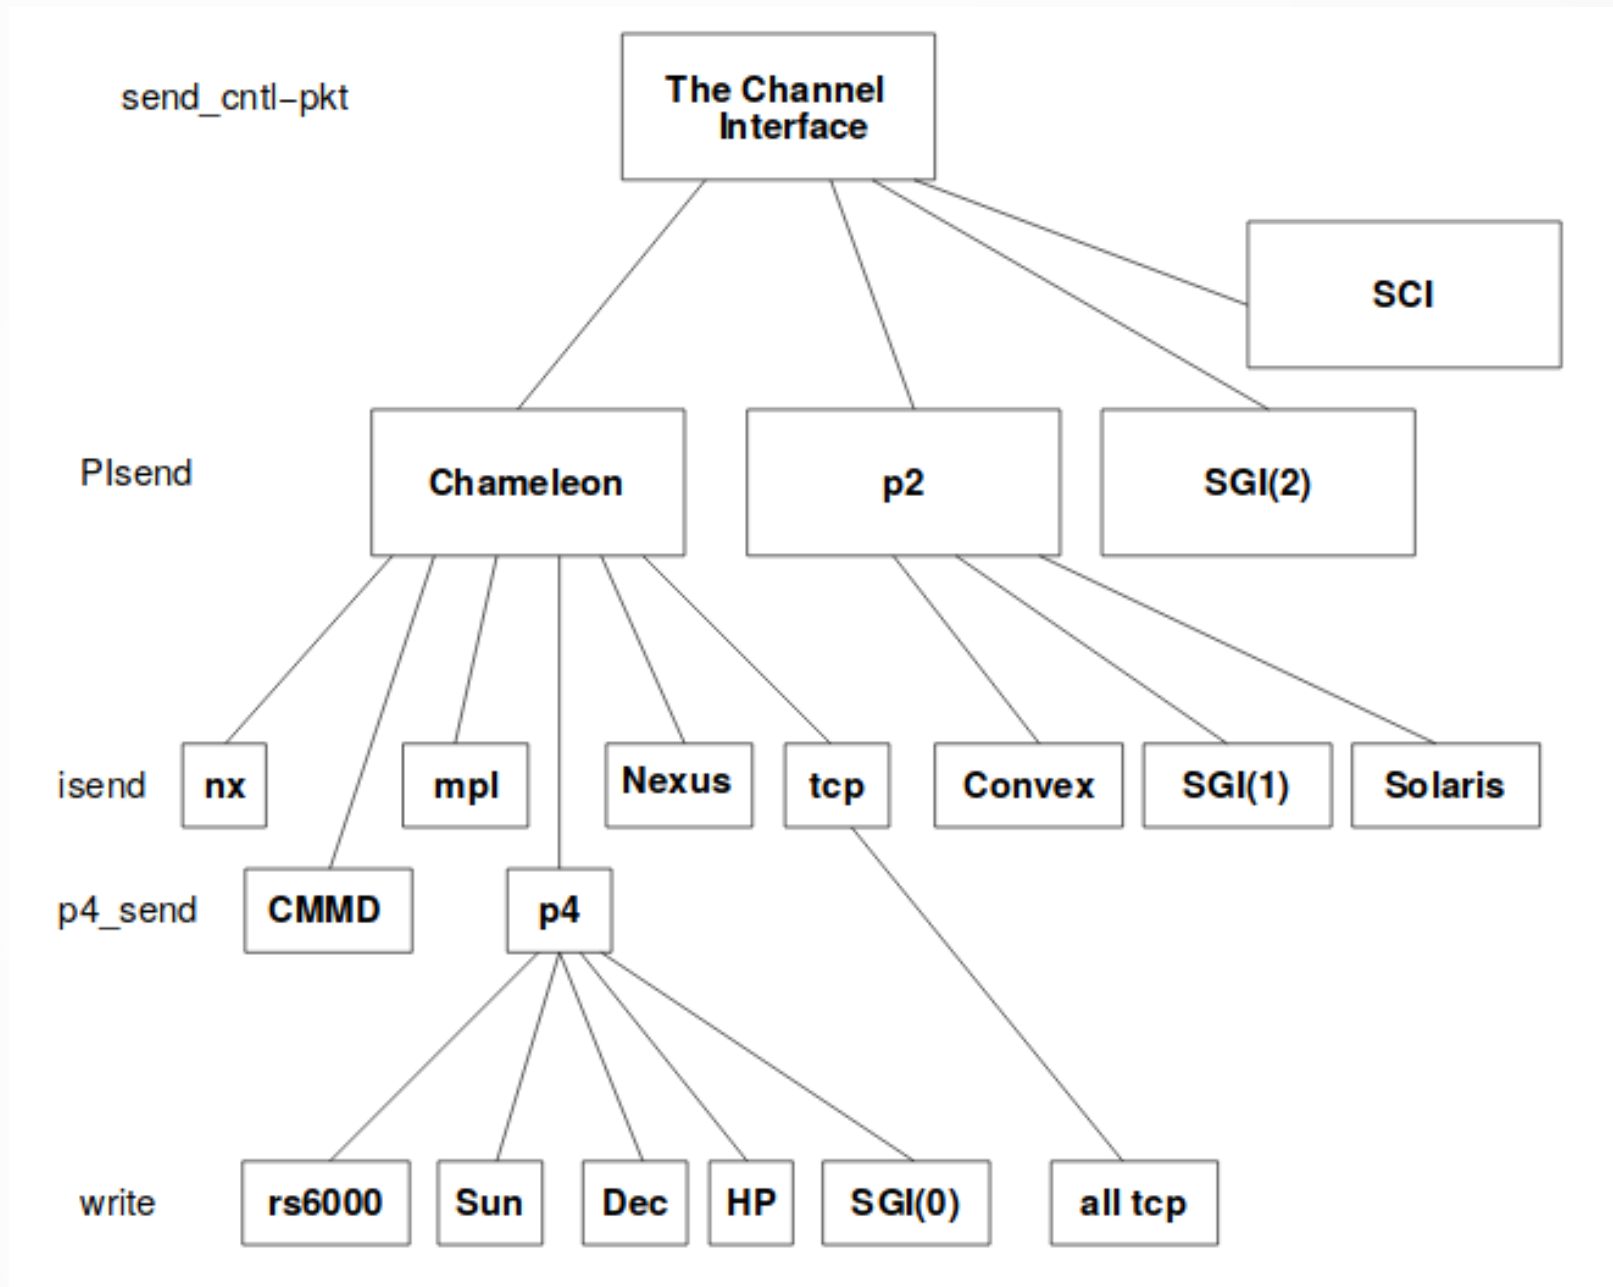
- Data exchange mechanisms
 - Eager
 - Data immediately sent to receiver
 - If receiver is not available to accept data, a buffer is allocated on receiving process to temporarily hold data

Channel Interface

- Data exchange mechanisms
 - Rendezvous
 - Data sent only when requested from receiver
 - Receiver sends a request for data and a way for the data to be transferred
 - Robust, but less efficient

Channel Interface

- Data exchange mechanisms
 - Get
 - Data read directly from receiver
 - Similar to rendezvous, but likely uses a form of **memcpy** to directly grab data
 - Requires specialized hardware (shared memory)



Exercise

- Revisit (circular-solved.c) using collective communication
 - Try a Scatter / Gather implementation
 - Try a Reduce implementation

Acknowledgements

- A High-Performance, Portable Implementation of the MPI Message Passing Interface Standard
 - Gropp, Lusk, Doss, Skjellum
 - http://web.cse.ohio-state.edu/~panda.2/788/papers/3a_P567.pdf